

Group12_Project_Implementation

April 18, 2024

1 Project Title: Diabetes Prediction - SVM Random Forest

1.1 Group Number: 12

1.1.1 Group Members:

1. Kheminda D.A.J.P. - Registration Number: EG/2019/3636
2. Kumarasinghe K.A.G.T.V - Registration Number: EG/2019/3645

2 Diabetes Prediction System

2.1 Import Required Dependencies

```
[162]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from imblearn.over_sampling import SMOTE
```

2.2 Dataset

```
[163]: # load the diabetes dataset
diabetes_dataset = pd.read_csv('diabetes.csv')

# print first 5 data rows in the dataset
diabetes_dataset.head(5)
```

```
[163]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

```
[164]: # to check the number of columns and rows in the dataset
diabetes_dataset.shape
```

```
[164]: (768, 9)
```

```
[165]: # Getting the statistical measures of the data
diabetes_dataset.describe()
```

```
[165]:
```

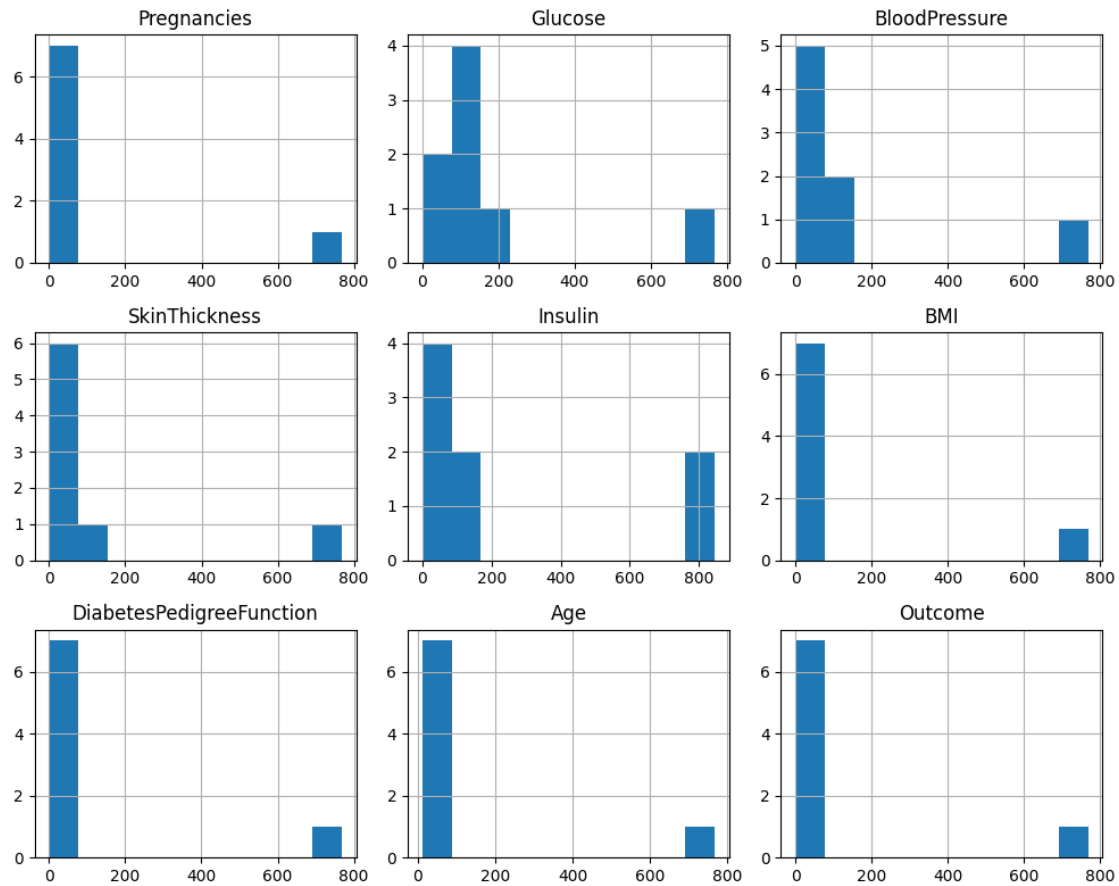
	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin \
count	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479
std	3.369578	31.972618	19.355807	15.952218	115.244002
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000
75%	6.000000	140.250000	80.000000	32.000000	127.250000
max	17.000000	199.000000	122.000000	99.000000	846.000000

	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000
mean	31.992578	0.471876	33.240885	0.348958
std	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.078000	21.000000	0.000000
25%	27.300000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000

2.2.1 Dataset columns

```
[166]: # data analysis and visualization to get a quick overview of the distribution
↳ of data
```

```
diabetes_dataset.describe().hist(figsize=(10, 8))
plt.tight_layout()
plt.show()
```




```
[167]: # to get idea about lables in Target Variable
diabetes_dataset['Outcome'].value_counts()
```

```
[167]: Outcome
0      500
1      268
Name: count, dtype: int64
```

```
[168]: #Visualization to get idea about lables in Target Variable
import seaborn as sns
sns.displot(diabetes_dataset['Outcome'])
```

```
[168]: <seaborn.axisgrid.FacetGrid at 0x78065cf8ecb0>
```



0 - Non diabetic 1 - diabetic

```
[169]: # Mean values of the columns
diabetes_dataset.groupby('Outcome').mean()
```

```
[169]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin \
Outcome					
0	3.298000	109.980000	68.184000	19.664000	68.792000
1	4.865672	141.257463	70.824627	22.164179	100.335821

	BMI	DiabetesPedigreeFunction	Age
Outcome			
0	30.304200	0.429734	31.190000
1	35.142537	0.550500	37.067164

2.3 Data Preprocessing

2.3.1 Handling null / missing values

```
[170]: #view dataframe summary
diabetes_dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies           768 non-null   int64
1   Glucose               768 non-null   int64
2   BloodPressure         768 non-null   int64
3   SkinThickness         768 non-null   int64
4   Insulin               768 non-null   int64
5   BMI                   768 non-null   float64
6   DiabetesPedigreeFunction 768 non-null   float64
7   Age                   768 non-null   int64
8   Outcome               768 non-null   int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

```
[171]: #Checking the null values
diabetes_dataset.isnull().sum()
```

```
[171]: Pregnancies           0
Glucose                 0
BloodPressure           0
SkinThickness           0
Insulin                 0
BMI                     0
DiabetesPedigreeFunction 0
Age                     0
Outcome                 0
dtype: int64
```

```
[172]: #Visualizing Missing data
!pip install missingno
import missingno as msno
msno.matrix(diabetes_dataset)
plt.xlabel("Missing Value Matrix")
plt.ylabel("Count")
```

```
Requirement already satisfied: missingno in /usr/local/lib/python3.10/dist-packages (0.5.2)
```


```
Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages
```

```

(from missingno) (1.25.2)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist-
packages (from missingno) (3.7.1)
Requirement already satisfied: scipy in /usr/local/lib/python3.10/dist-packages
(from missingno) (1.11.4)
Requirement already satisfied: seaborn in /usr/local/lib/python3.10/dist-
packages (from missingno) (0.13.1)
Requirement already satisfied: contourpy>=1.0.1 in
/usr/local/lib/python3.10/dist-packages (from matplotlib->missingno) (1.2.1)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-
packages (from matplotlib->missingno) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in
/usr/local/lib/python3.10/dist-packages (from matplotlib->missingno) (4.51.0)
Requirement already satisfied: kiwisolver>=1.0.1 in
/usr/local/lib/python3.10/dist-packages (from matplotlib->missingno) (1.4.5)
Requirement already satisfied: packaging>=20.0 in
/usr/local/lib/python3.10/dist-packages (from matplotlib->missingno) (24.0)
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/dist-
packages (from matplotlib->missingno) (9.4.0)
Requirement already satisfied: pyparsing>=2.3.1 in
/usr/local/lib/python3.10/dist-packages (from matplotlib->missingno) (3.1.2)
Requirement already satisfied: python-dateutil>=2.7 in
/usr/local/lib/python3.10/dist-packages (from matplotlib->missingno) (2.8.2)
Requirement already satisfied: pandas>=1.2 in /usr/local/lib/python3.10/dist-
packages (from seaborn->missingno) (2.0.3)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-
packages (from pandas>=1.2->seaborn->missingno) (2023.4)
Requirement already satisfied: tzdata>=2022.1 in /usr/local/lib/python3.10/dist-
packages (from pandas>=1.2->seaborn->missingno) (2024.1)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-
packages (from python-dateutil>=2.7->matplotlib->missingno) (1.16.0)

```

```
[172]: Text(0, 0.5, 'Count')
```



So, there is no any missing values in the dataset

2.3.2 Handling duplicated values

```
[173]: # Check for duplicated values
duplicated_rows = diabetes_dataset.duplicated()

# Display rows with duplicated values
duplicated_data = diabetes_dataset[duplicated_rows]
print("Duplicated Rows:")
print(duplicated_data)
```

Duplicated Rows:

Empty DataFrame

Columns: [Pregnancies, Glucose, BloodPressure, SkinThickness, Insulin, BMI, DiabetesPedigreeFunction, Age, Outcome]

Index: []

So, there is no any duplicated data

2.3.3 Handling Outliers

Detect outliers using Box-Plot

```
[174]: # Get list of features
features = diabetes_dataset.columns.tolist()

# Calculate the number of rows and columns for subplots
num_rows = 3
num_cols = 3
```

```

# Create subplots
fig, axes = plt.subplots(num_rows, num_cols, figsize=(15, 15))

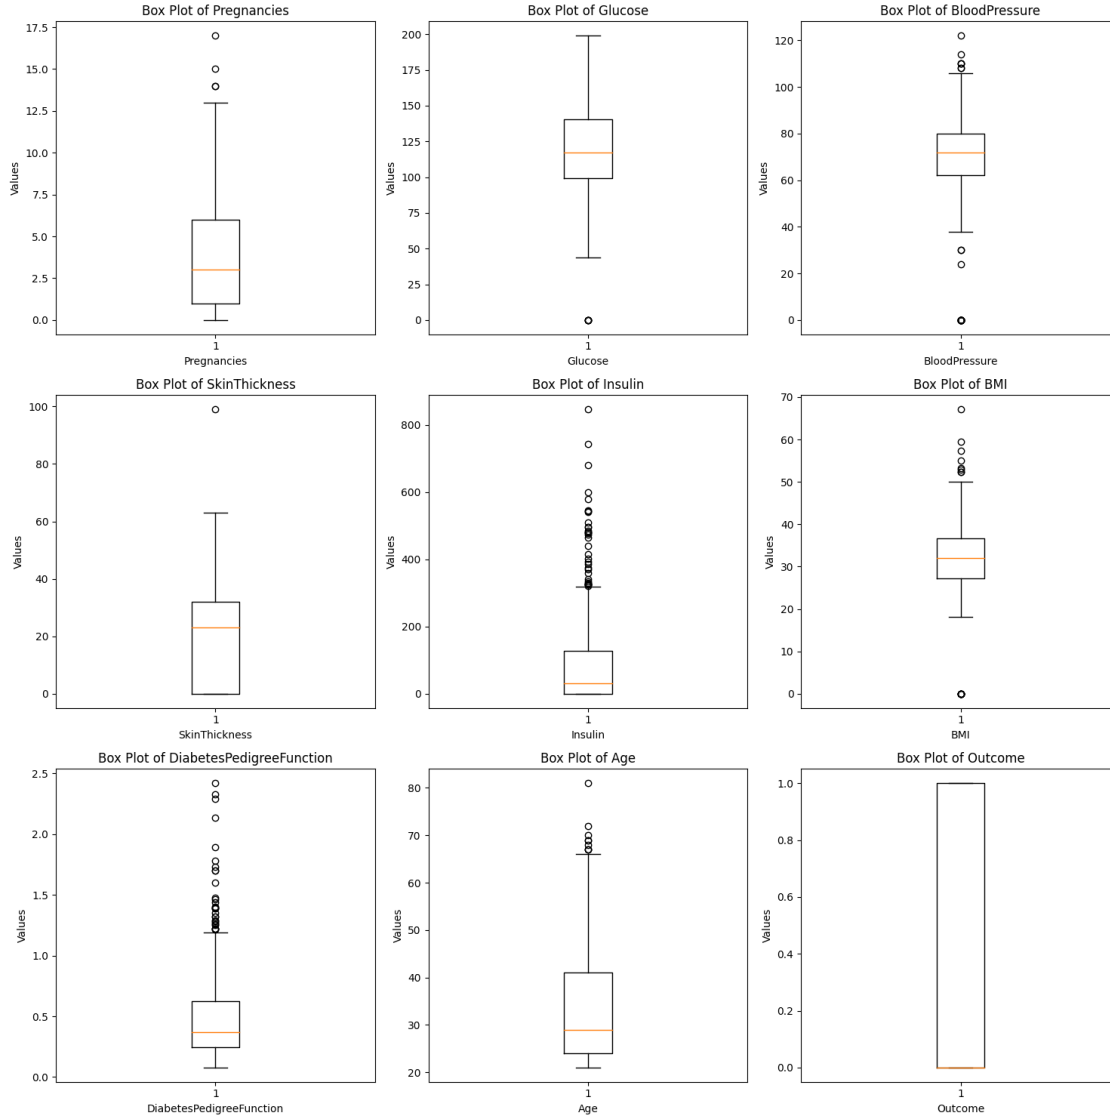
# Flatten the axes array
axes = axes.flatten()

# Create separate box plots for each feature
for i, feature in enumerate(features):
    ax = axes[i]
    ax.boxplot(diabetes_dataset[feature])
    ax.set_title('Box Plot of {}'.format(feature))
    ax.set_xlabel('{}'.format(feature))
    ax.set_ylabel('Values')

# Hide any extra subplots if the number of features is less than 9
for j in range(len(features), num_rows * num_cols):
    fig.delaxes(axes[j])

plt.tight_layout()
plt.show()

```

Treat to the outliers In our analysis to predict diabetes, we have chosen not to treat outliers in the dataset. Since outliers may contain valuable information relevant to the prediction task, such as DiabetesPedigreeFunction or insulin levels that could be indicative of diabetes, we have opted to retain them in our analysis. By keeping outliers, we aim to capture the full range of variability in the data, which may contribute to the accuracy and robustness of our predictive models.

2.3.4 Train Test Split

```
[175]: # Separate the data and labels
x = diabetes_dataset.drop('Outcome',axis=1)
y = diabetes_dataset['Outcome']
```

```
# print x and y
print(x)
print(y)
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	
..	
763	10	101	76	48	180	32.9	
764	2	122	70	27	0	36.8	
765	5	121	72	23	112	26.2	
766	1	126	60	0	0	30.1	
767	1	93	70	31	0	30.4	

	DiabetesPedigreeFunction	Age
0	0.627	50
1	0.351	31
2	0.672	32
3	0.167	21
4	2.288	33
..
763	0.171	63
764	0.340	27
765	0.245	30
766	0.349	47
767	0.315	23

[768 rows x 8 columns]

0	1
1	0
2	1
3	0
4	1
..	
763	0
764	0
765	0
766	1
767	0

Name: Outcome, Length: 768, dtype: int64

2.3.5 Treat for Imbalance dataset

```
[176]: # check the count of the target variable
y.value_counts()
```

```
[176]: Outcome
0      500
1      268
Name: count, dtype: int64
```

```
[177]: # Apply SMOTE
smote = SMOTE()
x_balanced, y_balanced = smote.fit_resample(x, y)

# Check the balanced target variable
print(pd.Series(y_balanced).value_counts())

# Check the x_balanced
print(x_balanced)
```

```
Outcome
1      500
0      500
Name: count, dtype: int64
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI \
0	6	148	72	35	0	33.600000
1	1	85	66	29	0	26.600000
2	8	183	64	0	0	23.300000
3	1	89	66	23	94	28.100000
4	0	137	40	35	168	43.100000
..
995	3	157	69	29	327	35.231892
996	6	146	89	25	479	31.110435
997	7	163	88	0	0	30.727504
998	3	173	82	24	160	31.365758
999	7	153	77	32	0	32.546891

	DiabetesPedigreeFunction	Age
0	0.627000	50
1	0.351000	31
2	0.672000	32
3	0.167000	21
4	2.288000	33
..
995	0.332693	34
996	0.145426	40
997	0.439279	53
998	1.007771	51

999 0.465523 45

[1000 rows x 8 columns]

```
[178]: diabetes_indicators = pd.DataFrame(x_balanced, columns=x.columns)
diabetes_indicators['Outcome'] = y_balanced

# Check the class distribution after applying SMOTE
diabetes_counts = diabetes_indicators['Outcome'].value_counts()


# Print the counts
for value, count in diabetes_counts.items():
    print(f"Count {value}: {count}")
```

Count 1: 500

Count 0: 500

```
[179]: #Visualization to get idea about lables in Target Variable after Smote
import seaborn as sns
sns.displot(diabetes_indicators['Outcome'])
```

```
[179]: <seaborn.axisgrid.FacetGrid at 0x78065ce8a560>
```



```
[180]: diabetes_indicators.describe()
```

```
[180]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin \
count	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000
mean	4.045000	126.018000	69.773000	20.738000	84.793000
std	3.369619	32.829917	19.420469	16.133308	123.387766
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1.000000	102.750000	64.000000	0.000000	0.000000
50%	3.000000	122.000000	72.000000	24.000000	0.000000
75%	6.000000	148.000000	80.000000	33.000000	136.250000
max	17.000000	199.000000	122.000000	99.000000	846.000000

	BMI	DiabetesPedigreeFunction	Age	Outcome
count	1000.000000	1000.000000	1000.000000	1000.000000
mean	32.736158	0.489805	34.292000	0.500000
std	7.478839	0.323174	11.587709	0.500250
min	0.000000	0.078000	21.000000	0.000000
25%	28.300000	0.256000	25.000000	0.000000

50%	32.800000	0.398500	31.000000	0.500000
75%	37.056576	0.647398	42.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000

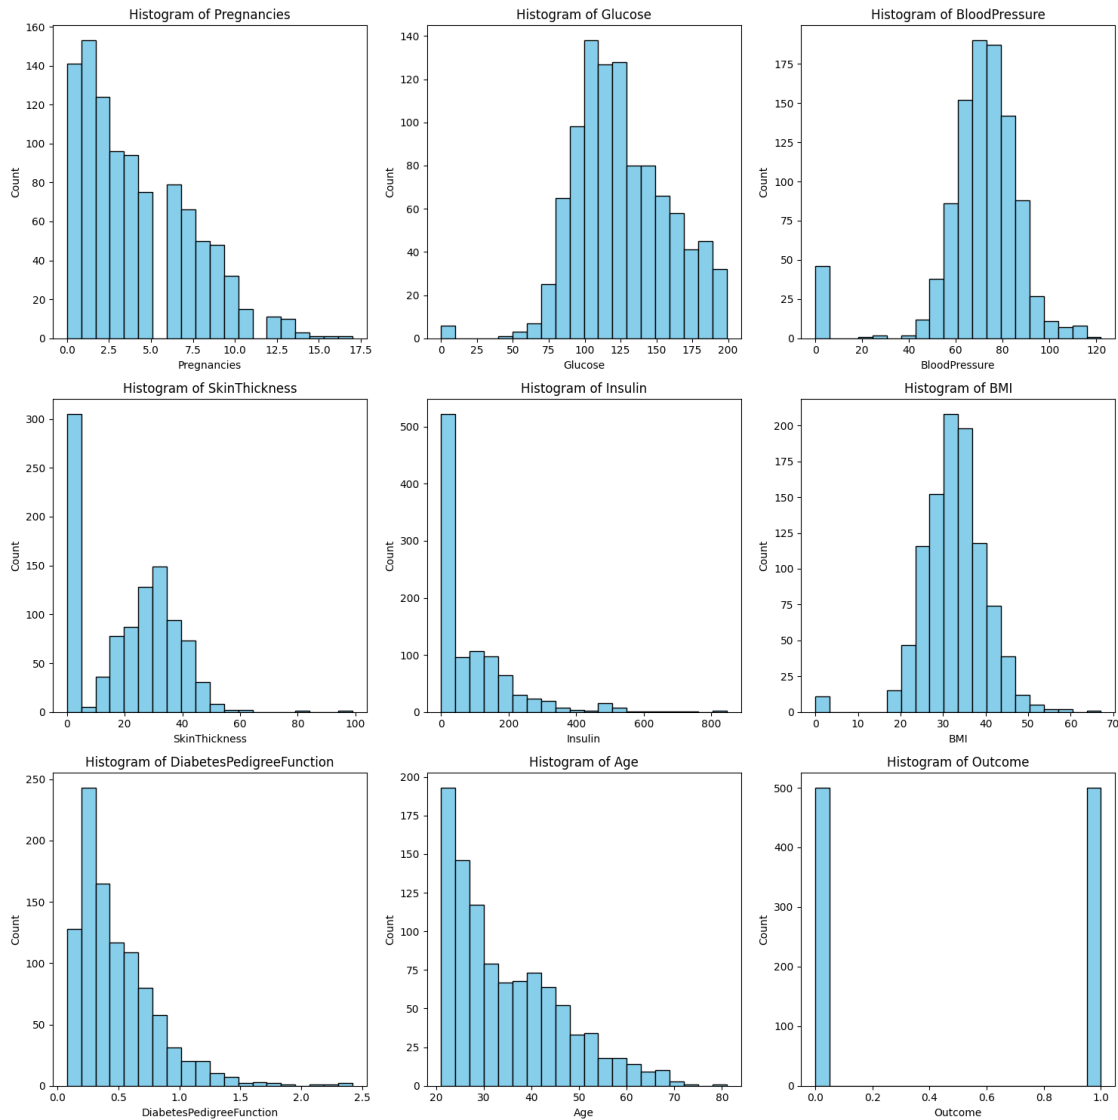
```
[181]: feature_names = diabetes_indicators.columns

# Create subplots for each feature
num_features = len(feature_names)
num_cols = 3 # Number of columns in the subplot grid
num_rows = (num_features + num_cols - 1) // num_cols # Number of rows needed

plt.figure(figsize=(15, 5*num_rows))

for i, feature in enumerate(feature_names, 1):
    plt.subplot(num_rows, num_cols, i)
    plt.hist(diabetes_indicators[feature], bins=20, color='skyblue',
             edgecolor='black')
    plt.xlabel(feature)
    plt.ylabel('Count')
    plt.title(f'Histogram of {feature}')


plt.tight_layout()
plt.show()
```



2.3.6 Correlation Matrix

```
[182]: # Plot the correlation matrix as a heatmap
correlation_matrix = diabetes_indicators.corr()

# Create a heatmap of the correlation matrix
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f",
            linewidths=.5)
plt.title('Correlation Matrix Heatmap with Numerical Values')
plt.show()
```



```
[183]: # separating the data and labels
x_balanced = diabetes_indicators.drop(columns = 'Outcome', axis=1)
y_balanced = diabetes_indicators['Outcome']
```

##Feature Scaling

2.3.7 Data Standadization

```
[184]: scaler = StandardScaler()
scaler.fit(x_balanced)
```

```
[184]: StandardScaler()
```

```
[185]: standardized_data = scaler.transform(x_balanced)

# print the data
```



```
standardized_data
```

```
[185]: array([[ 0.58047467,  0.66990724,  0.1147302 , ...,  0.11556261,
           0.4247363 ,  1.3562525 ],
          [-0.90411528, -1.25003435, -0.19437676, ..., -0.82087991,
           -0.42972118, -0.28423626],
          [ 1.17431065,  1.73654145, -0.29741242, ..., -1.26234567,
           0.56405002, -0.19789475],
          ...,
          [ 0.87739266,  1.12703618,  0.93901543, ..., -0.26871267,
           -0.15642186,  1.61527704],
          [-0.3102793 ,  1.43178882,  0.62990847, ..., -0.18332873,
           1.60355106,  1.44259401],
          [ 0.87739266,  0.82228355,  0.37231934, ..., -0.0253197 ,
           -0.07517313,  0.92454493]])
```

```
[186]: # again assign the standardized data into X variable
X = standardized_data

Y = y_balanced
```

2.3.8 Divide the data into two parts

```
[187]: x_train, x_test, y_train, y_test = train_test_split(X,Y, test_size = 0.2,
↳stratify =Y , random_state = 2)
```

```
[188]: # check the split

print(f'Original X : {X.shape}')
print(f'Original X : {Y.shape}')
print(f'Traing X   : {x_train.shape}')
print(f'Testing X  : {x_test.shape}')
print(f'Traing X   : {y_train.shape}')
print(f'Testing X  : {y_test.shape}')
```

```
Original X : (1000, 8)
Original X : (1000,)
Traing X   : (800, 8)
Testing X  : (200, 8)
Traing X   : (800,)
Testing X  : (200,)
```

2.4 Training the model

2.4.1 Using SVC (Support Vector Classifier)

```
[189]: # import the svc library
from sklearn import svm

# Train the model
model_svc = svm.SVC(kernel = 'linear')
model_svc.fit(x_train,y_train)
```

```
[189]: SVC(kernel='linear')
```

2.4.2 Using Random Forest

```
[190]: # import the svc library
from sklearn.ensemble import RandomForestClassifier

max_depth_value=5

# Train the model
model_RF = RandomForestClassifier(n_estimators=500,max_depth=max_depth_value)
model_RF.fit(x_train,y_train)
```

```
[190]: RandomForestClassifier(max_depth=5, n_estimators=500)
```

2.5 Model evaluation

2.5.1 Using accuracy score

```
[191]: # for svc model

# for train data
pred_svc_train = model_svc.predict(x_train)
accuracyScore_svc_train = accuracy_score(pred_svc_train,y_train)
print('Accuracy score of the train data using SVC          : ',
      ↪accuracyScore_svc_train)

# for test data
pred_svc = model_svc.predict(x_test)
accuracyScore_svc = accuracy_score(pred_svc,y_test)
print('Accuracy score of the test data using SVC          : ',
      ↪accuracyScore_svc)
```

```

# for Fandom Forest model

# for train data
pred_RF_train = model_RF.predict(x_train)
accuracyScore_RF_train = accuracy_score(pred_RF_train,y_train)
print('Accuracy score of the train data using RF : ', accuracyScore_RF_train)

# for test data
pred_RF = model_RF.predict(x_test)
accuracyScore_RF = accuracy_score(pred_RF,y_test)
print('Accuracy score of the test data using RF : ', accuracyScore_RF)

```

```

Accuracy score of the train data using SVC      : 0.7625
Accuracy score of the test data using SVC      : 0.765
Accuracy score of the train data using RF      : 0.85625
Accuracy score of the test data using RF      : 0.805

```

2.5.2 confusion matrix, classification report, and accuracy score to assess its performance.

```

[192]: from sklearn.metrics import classification_report, confusion_matrix, \
        accuracy_score
print('Metrics for Support Vector Machine')
print('-----')
print("Confusion Matrix :")
print(confusion_matrix(y_test,pred_svc))
print('-----')
print("Classification Report :")
print(classification_report(y_test,pred_svc))
print('-----')
print("accuracy score :",accuracy_score(y_test, pred_svc))

```

Metrics for Support Vector Machine

Confusion Matrix :

```

[[77 23]
 [24 76]]

```

Classification Report :

	precision	recall	f1-score	support
0	0.76	0.77	0.77	100
1	0.77	0.76	0.76	100
accuracy			0.77	200
macro avg	0.77	0.77	0.76	200
weighted avg	0.77	0.77	0.76	200

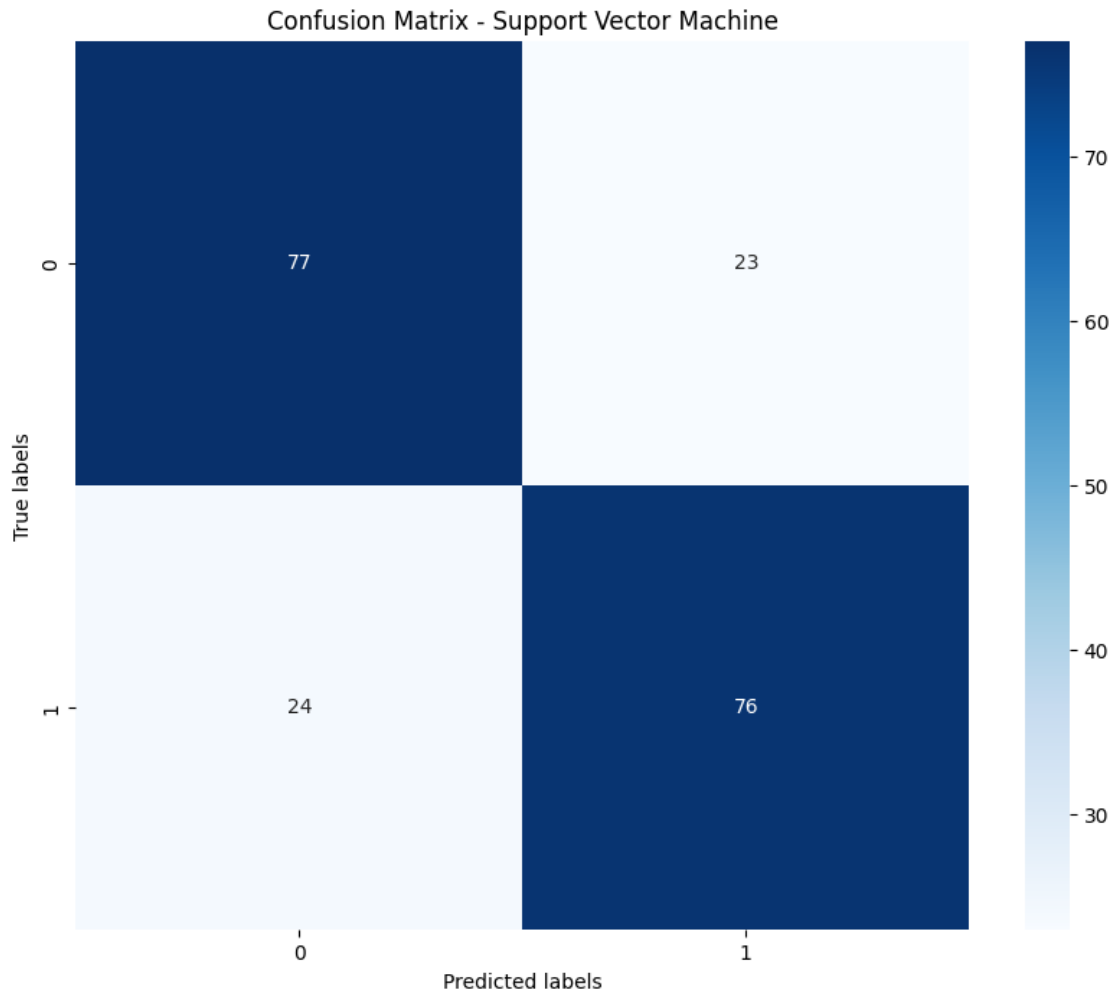
accuracy score : 0.765

```
[193]: import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix

# Get confusion matrix
cm = confusion_matrix(y_test, pred_svc)

# Define the unique class labels
classes = [0, 1] # Replace [0, 1] with your actual unique class labels

# Plot confusion matrix as heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(cm, annot=True, cmap='Blues', fmt='g',
            xticklabels=classes, yticklabels=classes)
plt.xlabel('Predicted labels')
plt.ylabel('True labels')
plt.title('Confusion Matrix - Support Vector Machine')
plt.show()
```



```
[194]: from sklearn.metrics import classification_report, confusion_matrix, \
        accuracy_score
print('Metrics for Random Forest')
print('-----')
print("Confusion Matrix :")
print(confusion_matrix(y_test,pred_RF))
print('-----')
print("Classification Report :")
print(classification_report(y_test,pred_RF))
print('-----')
print("accuracy score :",accuracy_score(y_test, pred_RF))
```

Metrics for Random Forest

Confusion Matrix :

[[77 23]

[16 84]]

Classification Report :

	precision	recall	f1-score	support
0	0.83	0.77	0.80	100
1	0.79	0.84	0.81	100
accuracy			0.81	200
macro avg	0.81	0.80	0.80	200
weighted avg	0.81	0.81	0.80	200


accuracy score : 0.805

```
[195]: import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix

# Get confusion matrix
cm = confusion_matrix(y_test, pred_RF)

# Define the unique class labels
classes = [0, 1] # Replace [0, 1] with your actual unique class labels

# Plot confusion matrix as heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(cm, annot=True, cmap='Blues', fmt='g',
            xticklabels=classes, yticklabels=classes)
plt.xlabel('Predicted labels')
plt.ylabel('True labels')
plt.title('Confusion Matrix - Random Forest')
plt.show()
```



2.6 Making the predictive system

```
[196]: input_data = (5,166,72,19,175,25.8,0.587,51)

# changing the input_data to numpy array
input_data_as_numpy_array = np.asarray(input_data)

# reshape the array as we are predicting for one instance
input_data_reshaped = input_data_as_numpy_array.reshape(1,-1)
```

2.6.1 Predict the diabetes using SVM

```
[197]: # Prediction using SVM classifier
svm_prediction = model_svc.predict(input_data_reshaped)
print("SVM Classifier Prediction:", svm_prediction[0])
```

```

if (svm_prediction[0] == 0):
    print('SVM : The person is not diabetic')
else:
    print('SVM: The person is diabetic')

```

SVM Classifier Prediction: 1
SVM: The person is diabetic

2.6.2 Predict the diabetes using RF

```

[198]: rf_prediction = model_RF.predict(input_data_reshaped)
print("Random Forest Classifier Prediction:", rf_prediction[0])

# Output prediction interpretation
if rf_prediction[0] == 0:
    print('Random Forest: The person is not diabetic')
else:
    print('Random Forest: The person is diabetic')

```

Random Forest Classifier Prediction: 1
Random Forest: The person is diabetic

2.6.3 Saving the trained SVM model & Random Forest model

```

[199]: import pickle

# Save the svm trained model using pickle
filename = 'DiabetesPrediction_svm_model.sav'
pickle.dump(model_svc, open(filename, 'wb'))

# Save the RF trained model using pickle
filename = 'DiabetesPrediction_rf_model.sav'
pickle.dump(model_RF, open(filename, 'wb'))

```

2.6.4 Loading the Save SVM model and Testing

```

[200]: load_model_svm = pickle.load(open('DiabetesPrediction_svm_model.sav', 'rb'))

```

```

[201]: # Prediction using SVM classifier
svm_prediction = load_model_svm.predict(input_data_reshaped)
print("SVM Classifier Prediction:", svm_prediction[0])

if (svm_prediction[0] == 0):
    print('SVM : The person is not diabetic')
else:
    print('SVM: The person is diabetic')

```



```
SVM Classifier Prediction: 1  
SVM: The person is diabetic
```

2.6.5 Loading the Save RF model and Testing

```
[202]: load_model_RF = pickle.load(open('DiabetesPrediction_rf_model.sav','rb'))
```

```
[203]: # Prediction using SVM classifier  
svm_prediction = load_model_RF.predict(input_data_reshaped)  
print("SVM Classifier Prediction:", svm_prediction[0])  
  
if (svm_prediction[0] == 0):  
    print('SVM : The person is not diabetic')  
else:  
    print('SVM: The person is diabetic')
```

```
SVM Classifier Prediction: 1  
SVM: The person is diabetic
```